Assignment I

Hadoop Exercise

19APSE4314

Tharushi Wijethunga

Department of Software Engineering

First Attempt

Step 1

Prerequisites

Docker: Ensure Docker is installed and running on your system. You can check if Docker is installed by running:
bash
Copy code

```
docker --version
```

Output

Step 2

Step 1: Pull the Hadoop Docker Image

1. **Choose a Hadoop Docker Image**:
   - Many Docker images are available for Hadoop single-node setups. For this lab, we'll use the `bde2020/hadoop-namenode` image from the Big Data Europe repository on Docker Hub.

Pull the Docker image:

```
docker pull bde2020/hadoop-namenode:latest
```

Then it generates error

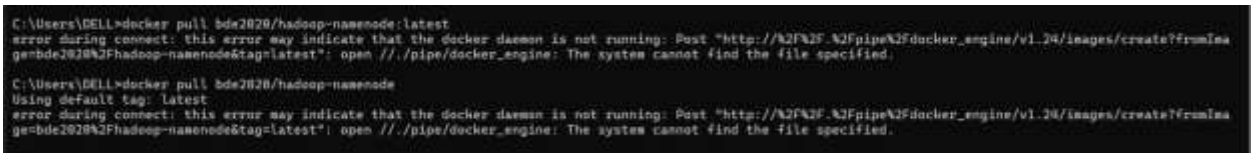

Start docker

```
3192219afd04: Pull complete
7127a1d8cced: Pull complete
883a09599980: Pull complete
77928a3e82af: Pull complete
92329e81aec4: Pull complete
f373218fec59: Pull complete
aa03513fe997: Pull complete
8b1800105b98: Pull complete
c3a84a3e49c8: Pull complete
a65640a64a76: Pull complete
a29cc756d786: Pull complete
abf352b16046: Pull complete
dddd5a009e99: Pull complete
Digest: sha256:fdf74118885132d646cf6f1263befc0919e1fb2ac5bd376c5366272fc261381e
Status: Downloaded newer image for bde2020/hadoop-namenode:latest
docker.io/bde2020/hadoop-namenode:latest

What's Next?
  View a summary of image vulnerabilities and recommendations → docker scout quickview bde2020/hadoop-namenode
```

Step 3

2. **Verify the Download**:

List Docker images to confirm the Hadoop image is downloaded:

```
docker images
```

```
C:\Users\DELL>docker images
REPOSITORY                  TAG       IMAGE ID       CREATED         SIZE
nginx                       latest    3b25b682ea82   6 weeks ago     192MB
demo                        v1        a92efe036fb2   8 months ago    187MB
<none>                      <none>    ef93dd23e683   8 months ago    187MB
demo                        v2        92b11f67b42b   9 months ago    187MB
redis                       alpine    3877664fc4fcf  10 months ago   41MB
alpine                      3.18.4    8ca4688f4f35   13 months ago   7.34MB
hello-world                 latest    d2c94e258dcb   18 months ago   13.3kB
demo                        latest    35c6c67ec35b   19 months ago   5.61MB
bde2020/hadoop-namenode     latest    b6383071a2119  4 years ago     1.37GB
```

Step 4

Step 2: Start the Hadoop Container

1. **Run the Container**:

Start a container with the necessary configurations to act as a single-node Hadoop cluster:

```
docker run -it --name hadoop-cluster -p 9870:9870 -p 8088:8088 -p
50070:50070 bde2020/hadoop-namenode:latest /bin/bash
```

```
C:\Users\DELL>docker run -it --name hadoop-cluster -p 9870:9870 -p 8088:8088 -p 50070:50070 bde2020/hadoop-namenode:latest /bin/bash
Configuring core
 - Setting fs.defaultFS=hdfs://42d338f08c67:8020
Configuring hdfs
 - Setting dfs.namenode.name.dir=file:///hadoop/dfs/name
Configuring yarn
Configuring httpfs
Configuring kms
Configuring mapred
Configuring for multihomed network
root@42d338f08c67:/#
```

Step 5

2. **Start Hadoop Services**:

Once inside the container's shell, start the Hadoop services:

```
start-all.sh
```

- This will initialize HDFS and YARN.

```
root@42d338f08c67:/# find / -name start-all.sh
/opt/hadoop-3.2.1/sbin/start-all.sh
root@42d338f08c67:/# /opt/hadoop-3.2.1/sbin/start-all.sh
Starting namenodes on [42d338f08c67]
ERROR: Attempting to operate on hdfs namenode as root
ERROR: but there is no HDFS_NAMENODE_USER defined. Aborting operation.
Starting datanodes
ERROR: Attempting to operate on hdfs datanode as root
ERROR: but there is no HDFS_DATANODE_USER defined. Aborting operation.
Starting secondary namenodes [42d338f08c67]
ERROR: Attempting to operate on hdfs secondarynamenode as root
```

From this moments onwards the above error remains same

Second Attempt

Then I have start from the beginning in the powershell

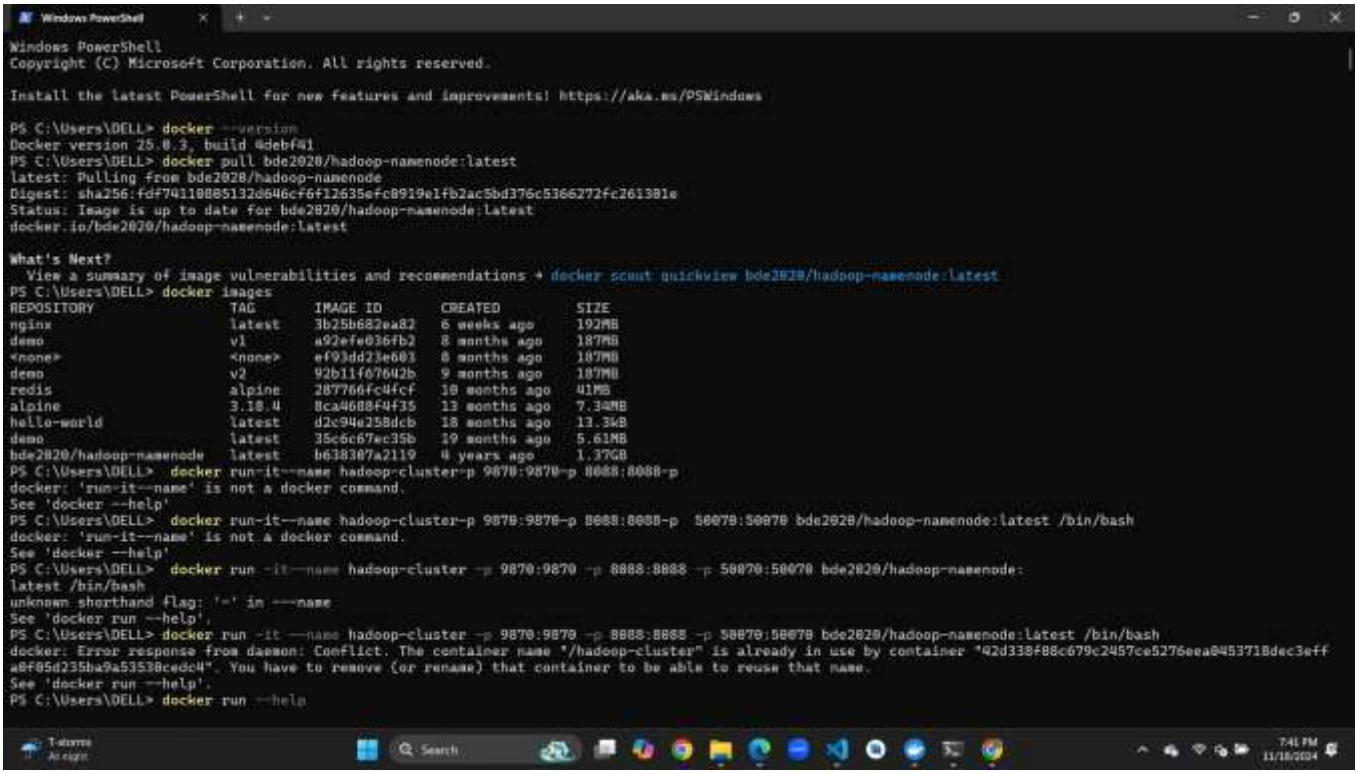

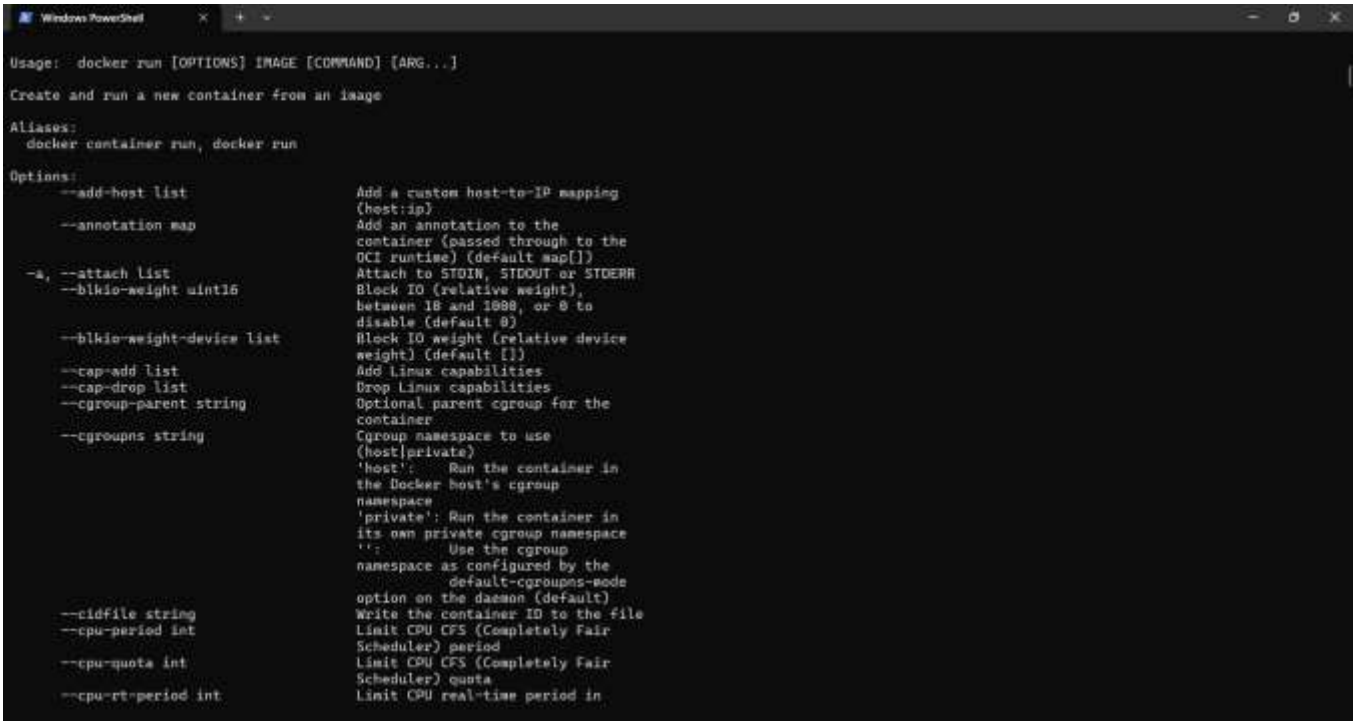

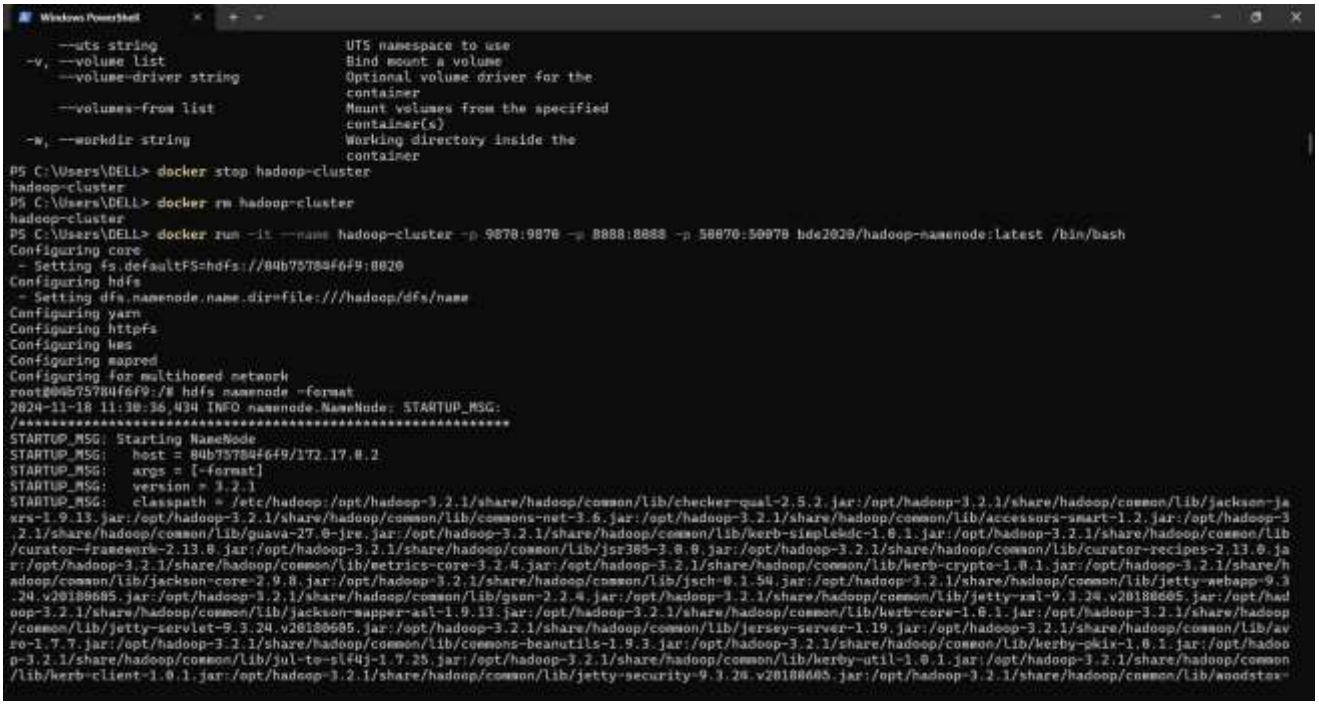
Windows PowerShell
--uts string UTS namespace to use
-v, --volume list Bind mount a volume
--volume-driver string Optional volume driver for the container
--volumes-from list Mount volumes from the specified container(s)
-w, --workdir string Working directory inside the container
PS C:\Users\DELL> docker stop hadoop-cluster
hadoop-cluster
PS C:\Users\DELL> docker rm hadoop-cluster
hadoop-cluster
PS C:\Users\DELL> docker run -it --name hadoop-cluster -p 9870:9870 -p 8088:8088 -p 50070:50070 bde2020/hadoop-namenode:latest /bin/bash
Configuring core
- Setting fs.defaultFS=hdfs://84b75784f6f9:8020
Configuring hdfs
- Setting dfs.namenode.name.dir=file:///hadoop/dfs/name
Configuring yarn
Configuring httpfs
Configuring kms
Configuring mapred
Configuring for multihomed network
root@84b75784f6f9:/# hdfs namenode -format
2024-11-18 11:30:36,434 INFO namenode.NameNode: STARTUP_MSG:
/**
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = 84b75784f6f9/172.17.0.2
STARTUP_MSG: args = [-format]
STARTUP_MSG: version = 3.2.1
STARTUP_MSG: classpath = /etc/hadoop:/opt/hadoop-3.2.1/share/hadoop/common/lib/checker-qual-2.5.2.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jackson-jaxrs-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-net-3.6.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/accessors-smart-1.2.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/guava-27.0-jre.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-simplekdc-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/curator-framework-2.13.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jsr305-3.0.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/curator-recipes-2.13.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/metrics-core-3.2.4.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-crypto-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jackson-core-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jsch-0.1.54.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-webapp-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/gson-2.2.4.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-xml-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jackson-mapper-asl-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-core-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jersey-server-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/avro-1.7.7.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-beanutils-1.9.3.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-pkix-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jul-to-slf4j-1.7.25.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-util-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-client-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-security-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/woodstox-


Windows PowerShell
hdfs/lib/jaxb-api-2.2.11.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/commons-math3-3.1.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jackson-annotations-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jetty-util-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/httpcore-4.4.10.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jackson-core-asl-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/commons-lang3-3.7.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/coomons-collections-3.2.2.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/javax.servlet-api-3.1.0.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/kerb-admin-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/commons-daemon-1.0.13.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/hadoop-auth-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/listenablefuture-9999.0-empty-to-avoid-conflict-with-guava.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/protobuf-java-2.5.0.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/kerb-identity-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/leveldbjni-all-1.8.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jersey-servlet-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/netty-all-4.0.52.Final.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jetty-io-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/okhttp-2.7.5.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/asm-5.0.4.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jersey-core-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jetty-server-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/stax2-api-3.1.4.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/re2j-1.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/jetty-util-ajax-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/lib/paranamer-2.3.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-client-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-client-3.2.1-tests.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-nfs-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-httpfs-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-rbf-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-native-client-3.2.1-tests.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-3.2.1-tests.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-rbf-3.2.1-tests.jar:/opt/hadoop-3.2.1/share/hadoop/hdfs/hadoop-hdfs-native-client-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/lib/hamcrest-core-1.3.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/lib/junit-4.11.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-app-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-nativetask-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-hs-plugins-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-jobclient-3.2.1-tests.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-shuffle-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-jobclient-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-common-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-uploader-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-core-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-client-hs-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/metrics-core-3.2.4.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/jersey-client-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/guice-servlet-4.0.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/bcpkix-jdk15on-1.60.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/jackson-module-jaxb-annotations-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/swagger-annotations-1.5.4.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/mssql-jdbc-6.2.1.jre7.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/java-util-1.9.0.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/snakeyaml-1.16.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/fst-2.50.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/json-io-2.5.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/geronimo-jcache_1.0_spec-1.0-alpha-1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/javax.inject-1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/objenesis-1.0.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/ehcache-3.3.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/HikariCP-java7-2.4.12.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/jackson-jaxrs-json-provider-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/jersey-guice-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/guice-4.0.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/jackson-jaxrs-base-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/aopalliance-1.0.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/lib/bcprov-jdk15on-1.60.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-client-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-services-api-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-submarine-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-registry-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-services-core-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-resourcemanager-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-applications-distributedshell-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-sharedcachemanager-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-applications-unmanaged-am-launcher-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-web-proxy-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-api-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-router-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-common-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-applicationhistoryservice-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-nodemanager-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-tests-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-common-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-timeline-pluginstorage-3.2.1.jar

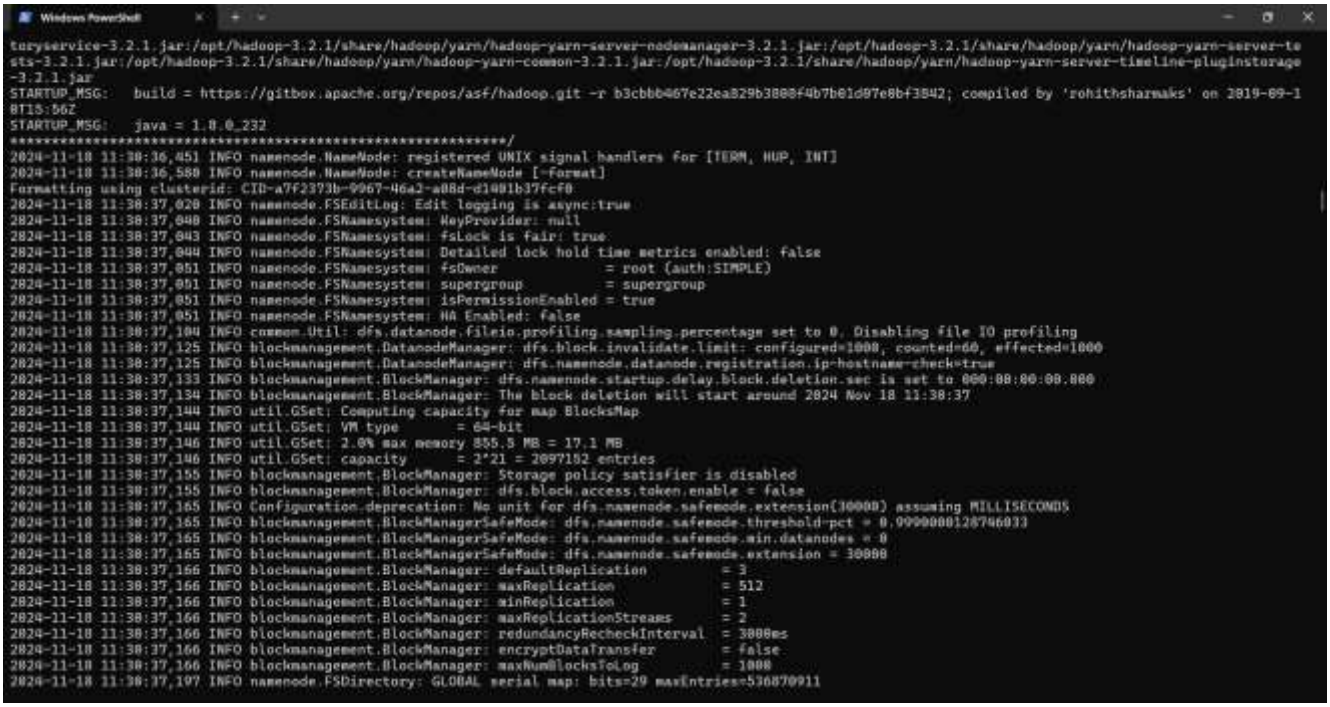
Windows PowerShell
toryservice-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-nodemanager-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-tests-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-common-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-timeline-pluginstorage-3.2.1.jar
STARTUP_MSG: build = https://gitbox.apache.org/repos/asf/hadoop.git -r b3cbbb467e22ea829b3808f4b7b01d07e0bf3842; compiled by 'rohithsharmaks' on 2019-09-10T15:56Z
STARTUP_MSG: java = 1.8.0_232
**/
2024-11-18 11:30:36,451 INFO namenode.NameNode: registered UNIX signal handlers for [TERM, HUP, INT]
2024-11-18 11:30:36,580 INFO namenode.NameNode: createNameNode [-format]
Formatting using clusterid: CID-a7f2373b-9967-46a2-a08d-d1401b37fcf0
2024-11-18 11:30:37,020 INFO namenode.FSEditLog: Edit logging is async:true
2024-11-18 11:30:37,040 INFO namenode.FSNamesystem: KeyProvider: null
2024-11-18 11:30:37,043 INFO namenode.FSNamesystem: fsLock is fair: true
2024-11-18 11:30:37,044 INFO namenode.FSNamesystem: Detailed lock hold time metrics enabled: false
2024-11-18 11:30:37,051 INFO namenode.FSNamesystem: fsOwner = root (auth:SIMPLE)
2024-11-18 11:30:37,051 INFO namenode.FSNamesystem: supergroup = supergroup
2024-11-18 11:30:37,051 INFO namenode.FSNamesystem: isPermissionEnabled = true
2024-11-18 11:30:37,051 INFO namenode.FSNamesystem: HA Enabled: false
2024-11-18 11:30:37,104 INFO common.Util: dfs.datanode.fileio.profiling.sampling.percentage set to 0. Disabling file IO profiling
2024-11-18 11:30:37,125 INFO blockmanagement.DatanodeManager: dfs.block.invalidate.limit: configured=1000, counted=60, effected=1000
2024-11-18 11:30:37,125 INFO blockmanagement.DatanodeManager: dfs.namenode.datanode.registration.ip-hostname-check=true
2024-11-18 11:30:37,133 INFO blockmanagement.BlockManager: dfs.namenode.startup.delay.block.deletion.sec is set to 000:00:00:00.000
2024-11-18 11:30:37,134 INFO blockmanagement.BlockManager: The block deletion will start around 2024 Nov 18 11:30:37
2024-11-18 11:30:37,144 INFO util.GSet: Computing capacity for map BlocksMap
2024-11-18 11:30:37,144 INFO util.GSet: VM type = 64-bit
2024-11-18 11:30:37,146 INFO util.GSet: 2.0% max memory 855.5 MB = 17.1 MB
2024-11-18 11:30:37,146 INFO util.GSet: capacity = 2^21 = 2097152 entries
2024-11-18 11:30:37,155 INFO blockmanagement.BlockManager: Storage policy satisfier is disabled
2024-11-18 11:30:37,155 INFO blockmanagement.BlockManager: dfs.block.access.token.enable = false
2024-11-18 11:30:37,165 INFO Configuration.deprecation: No unit for dfs.namenode.safemode.extension(30000) assuming MILLISECONDS
2024-11-18 11:30:37,165 INFO blockmanagement.BlockManagerSafeMode: dfs.namenode.safemode.threshold-pct = 0.9990000128746033
2024-11-18 11:30:37,165 INFO blockmanagement.BlockManagerSafeMode: dfs.namenode.safemode.min.datanodes = 0
2024-11-18 11:30:37,165 INFO blockmanagement.BlockManagerSafeMode: dfs.namenode.safemode.extension = 30000
2024-11-18 11:30:37,166 INFO blockmanagement.BlockManager: defaultReplication = 3
2024-11-18 11:30:37,166 INFO blockmanagement.BlockManager: maxReplication = 512
2024-11-18 11:30:37,166 INFO blockmanagement.BlockManager: minReplication = 1
2024-11-18 11:30:37,166 INFO blockmanagement.BlockManager: maxReplicationStreams = 2
2024-11-18 11:30:37,166 INFO blockmanagement.BlockManager: redundancyRecheckInterval = 3000ms
2024-11-18 11:30:37,166 INFO blockmanagement.BlockManager: encryptDataTransfer = false
2024-11-18 11:30:37,166 INFO blockmanagement.BlockManager: maxNumBlocksToLog = 1000
2024-11-18 11:30:37,197 INFO namenode.FSDirectory: GLOBAL serial map: bits=29 maxEntries=536870911

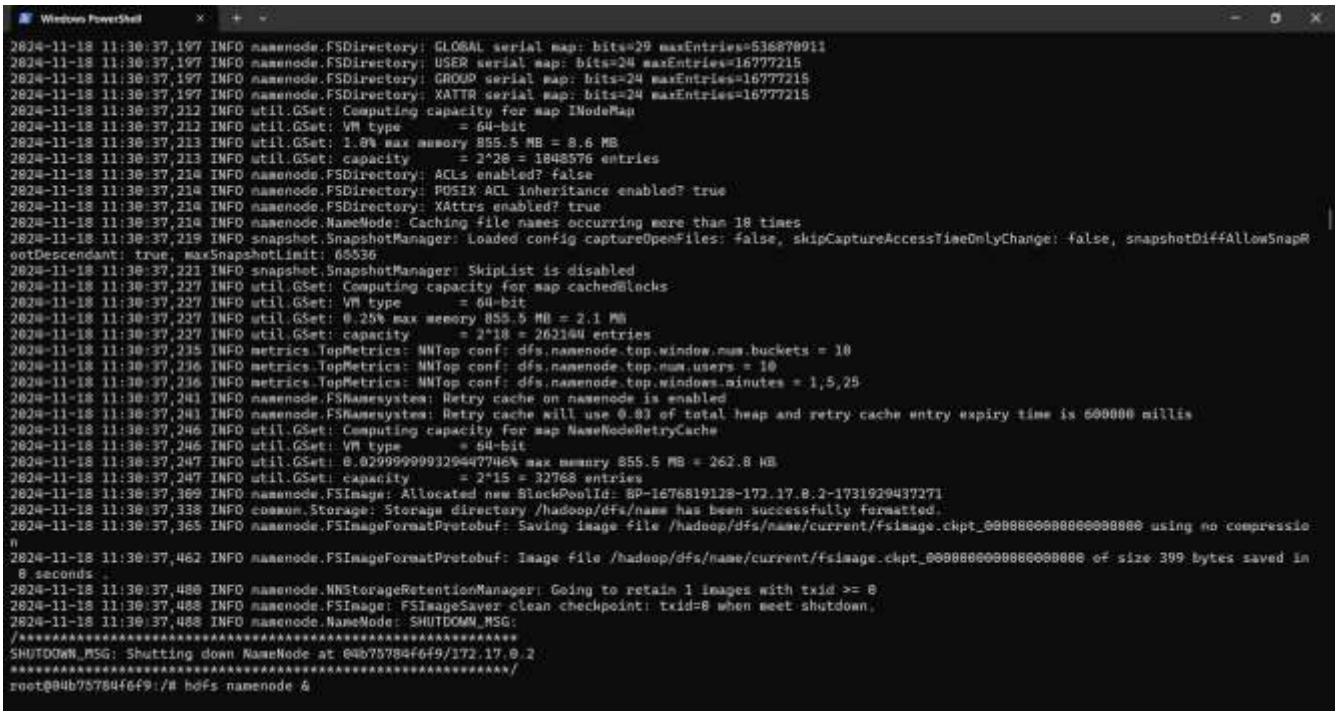
Windows PowerShell
2024-11-18 11:30:37,197 INFO namenode.FSDirectory: GLOBAL serial map: bits=29 maxEntries=536870911
2024-11-18 11:30:37,197 INFO namenode.FSDirectory: USER serial map: bits=24 maxEntries=16777215
2024-11-18 11:30:37,197 INFO namenode.FSDirectory: GROUP serial map: bits=24 maxEntries=16777215
2024-11-18 11:30:37,197 INFO namenode.FSDirectory: XATTR serial map: bits=24 maxEntries=16777215
2024-11-18 11:30:37,212 INFO util.GSet: Computing capacity for map INodeMap
2024-11-18 11:30:37,212 INFO util.GSet: VM type = 64-bit
2024-11-18 11:30:37,213 INFO util.GSet: 1.0% max memory 855.5 MB = 8.6 MB
2024-11-18 11:30:37,213 INFO util.GSet: capacity = 2^20 = 1048576 entries
2024-11-18 11:30:37,214 INFO namenode.FSDirectory: ACLs enabled? false
2024-11-18 11:30:37,214 INFO namenode.FSDirectory: POSIX ACL inheritance enabled? true
2024-11-18 11:30:37,214 INFO namenode.FSDirectory: XAttrs enabled? true
2024-11-18 11:30:37,214 INFO namenode.NameNode: Caching file names occurring more than 10 times
2024-11-18 11:30:37,219 INFO snapshot.SnapshotManager: Loaded config captureOpenFiles: false, skipCaptureAccessTimeOnlyChange: false, snapshotDiffAllowSnapRootDescendant: true, maxSnapshotLimit: 65536
2024-11-18 11:30:37,221 INFO snapshot.SnapshotManager: SkipList is disabled
2024-11-18 11:30:37,227 INFO util.GSet: Computing capacity for map cachedBlocks
2024-11-18 11:30:37,227 INFO util.GSet: VM type = 64-bit
2024-11-18 11:30:37,227 INFO util.GSet: 0.25% max memory 855.5 MB = 2.1 MB
2024-11-18 11:30:37,227 INFO util.GSet: capacity = 2^18 = 262144 entries
2024-11-18 11:30:37,235 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.window.num.buckets = 10
2024-11-18 11:30:37,236 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.num.users = 10
2024-11-18 11:30:37,236 INFO metrics.TopMetrics: NNTop conf: dfs.namenode.top.windows.minutes = 1,5,25
2024-11-18 11:30:37,241 INFO namenode.FSNamesystem: Retry cache on namenode is enabled
2024-11-18 11:30:37,241 INFO namenode.FSNamesystem: Retry cache will use 0.03 of total heap and retry cache entry expiry time is 600000 millis
2024-11-18 11:30:37,246 INFO util.GSet: Computing capacity for map NameNodeRetryCache
2024-11-18 11:30:37,246 INFO util.GSet: VM type = 64-bit
2024-11-18 11:30:37,247 INFO util.GSet: 0.029999999329447746% max memory 855.5 MB = 262.8 KB
2024-11-18 11:30:37,247 INFO util.GSet: capacity = 2^15 = 32768 entries
2024-11-18 11:30:37,309 INFO namenode.FSImage: Allocated new BlockPoolId: BP-1676819128-172.17.8.2-1731929437271
2024-11-18 11:30:37,338 INFO common.Storage: Storage directory /hadoop/dfs/name has been successfully formatted.
2024-11-18 11:30:37,365 INFO namenode.FSImageFormatProtobuf: Saving image file /hadoop/dfs/name/current/fsimage.ckpt_0000000000000000000 using no compression
2024-11-18 11:30:37,462 INFO namenode.FSImageFormatProtobuf: Image file /hadoop/dfs/name/current/fsimage.ckpt_0000000000000000000 of size 399 bytes saved in 0 seconds .
2024-11-18 11:30:37,480 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2024-11-18 11:30:37,488 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2024-11-18 11:30:37,488 INFO namenode.NameNode: SHUTDOWN_MSG:
/**
SHUTDOWN_MSG: Shutting down NameNode at 00b75784f6f9/172.17.0.2
**/
root@00b75784f6f9:/# hdfs namenode &

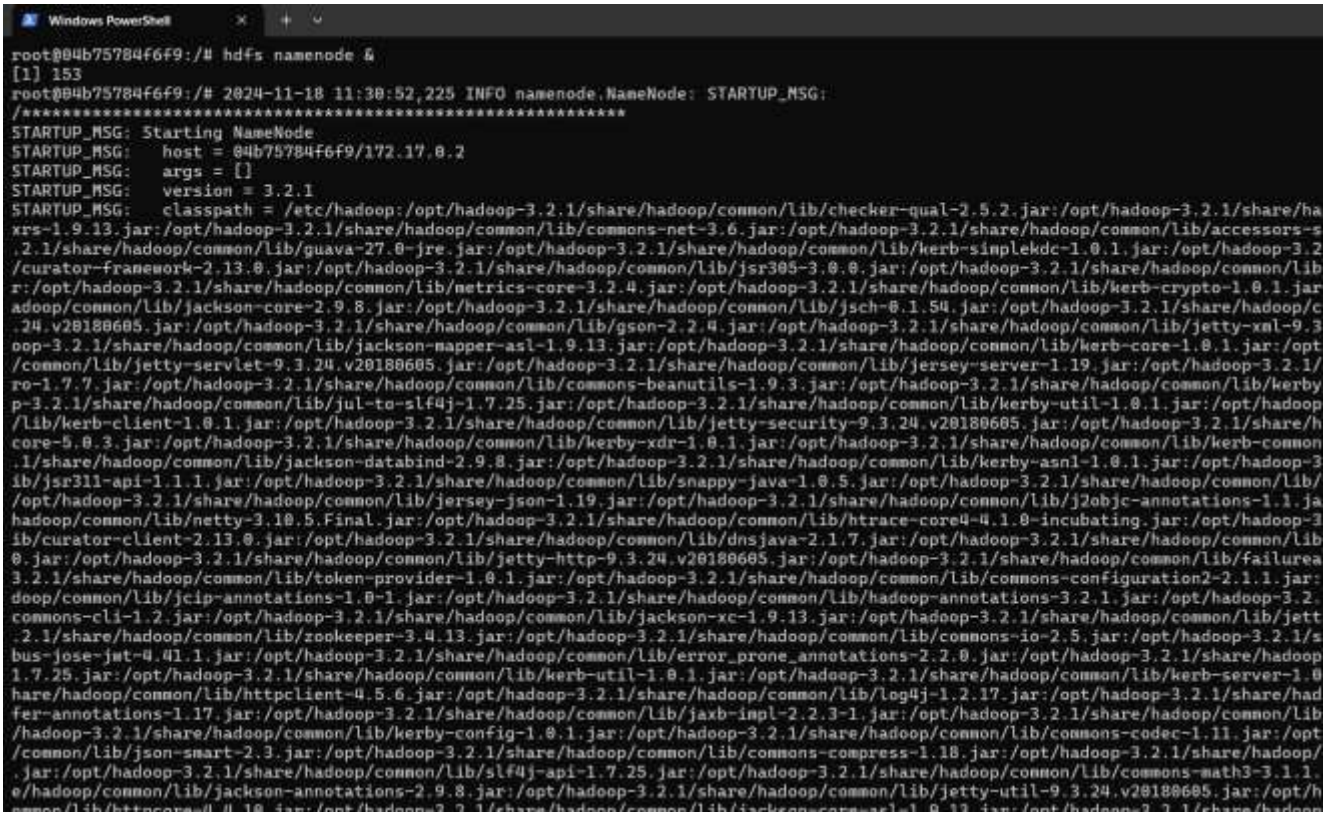
Windows PowerShell
root@04b75784f6f9:/# hdfs namenode &
[1] 153
root@04b75784f6f9:/# 2024-11-18 11:30:52,225 INFO namenode.NameNode: STARTUP_MSG:
/**
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = 04b75784f6f9/172.17.0.2
STARTUP_MSG: args = []
STARTUP_MSG: version = 3.2.1
STARTUP_MSG: classpath = /etc/hadoop:/opt/hadoop-3.2.1/share/hadoop/common/lib/checker-qual-2.5.2.jar:/opt/hadoop-3.2.1/share/ha
xrs-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-net-3.6.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/accessors-s
.2.1/share/hadoop/common/lib/guava-27.0-jre.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-simplekdc-1.0.1.jar:/opt/hadoop-3.2
/curator-framework-2.13.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jsr305-3.0.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib
r:/opt/hadoop-3.2.1/share/hadoop/common/lib/metrics-core-3.2.4.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-crypto-1.0.1.jar
adoop/common/lib/jackson-core-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jsch-0.1.54.jar:/opt/hadoop-3.2.1/share/hadoop/c
.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/gson-2.2.4.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-xml-9.3
oop-3.2.1/share/hadoop/common/lib/jackson-mapper-asl-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-core-1.0.1.jar:/opt
/common/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jersey-server-1.19.jar:/opt/hadoop-3.2.1/
ro-1.7.7.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-beanutils-1.9.3.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby
p-3.2.1/share/hadoop/common/lib/jul-to-slf4j-1.7.25.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-util-1.0.1.jar:/opt/hadoop
/lib/kerb-client-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-security-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/h
core-5.0.3.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-xdr-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-common
.1/share/hadoop/common/lib/jackson-databind-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-asn1-1.0.1.jar:/opt/hadoop-3
ib/jsr311-api-1.1.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/snappy-java-1.0.5.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/
/opt/hadoop-3.2.1/share/hadoop/common/lib/jersey-json-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/j2objc-annotations-1.1.ja
hadoop/common/lib/netty-3.10.5.Final.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/htrace-core4-4.1.0-incubating.jar:/opt/hadoop-3
ib/curator-client-2.13.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/dnsjava-2.1.7.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib
0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-http-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/failurea
3.2.1/share/hadoop/common/lib/token-provider-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-configuration2-2.1.1.jar:
doop/common/lib/jcip-annotations-1.0-1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/hadoop-annotations-3.2.1.jar:/opt/hadoop-3.2.
commons-cli-1.2.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jackson-xc-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jett
.2.1/share/hadoop/common/lib/zookeeper-3.4.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-io-2.5.jar:/opt/hadoop-3.2.1/s
bus-jose-jwt-4.41.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/error_prone_annotations-2.2.0.jar:/opt/hadoop-3.2.1/share/hadoop
1.7.25.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-util-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-server-1.0
hare/hadoop/common/lib/httpclient-4.5.6.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/log4j-1.2.17.jar:/opt/hadoop-3.2.1/share/had
fer-annotations-1.17.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jaxb-impl-2.2.3-1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib
/hadoop-3.2.1/share/hadoop/common/lib/kerby-config-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-codec-1.11.jar:/opt
/common/lib/json-smart-2.3.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-compress-1.18.jar:/opt/hadoop-3.2.1/share/hadoop/
.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/slf4j-api-1.7.25.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-math3-3.1.1.
e/hadoop/common/lib/jackson-annotations-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-util-9.3.24.v20180605.jar:/opt/h

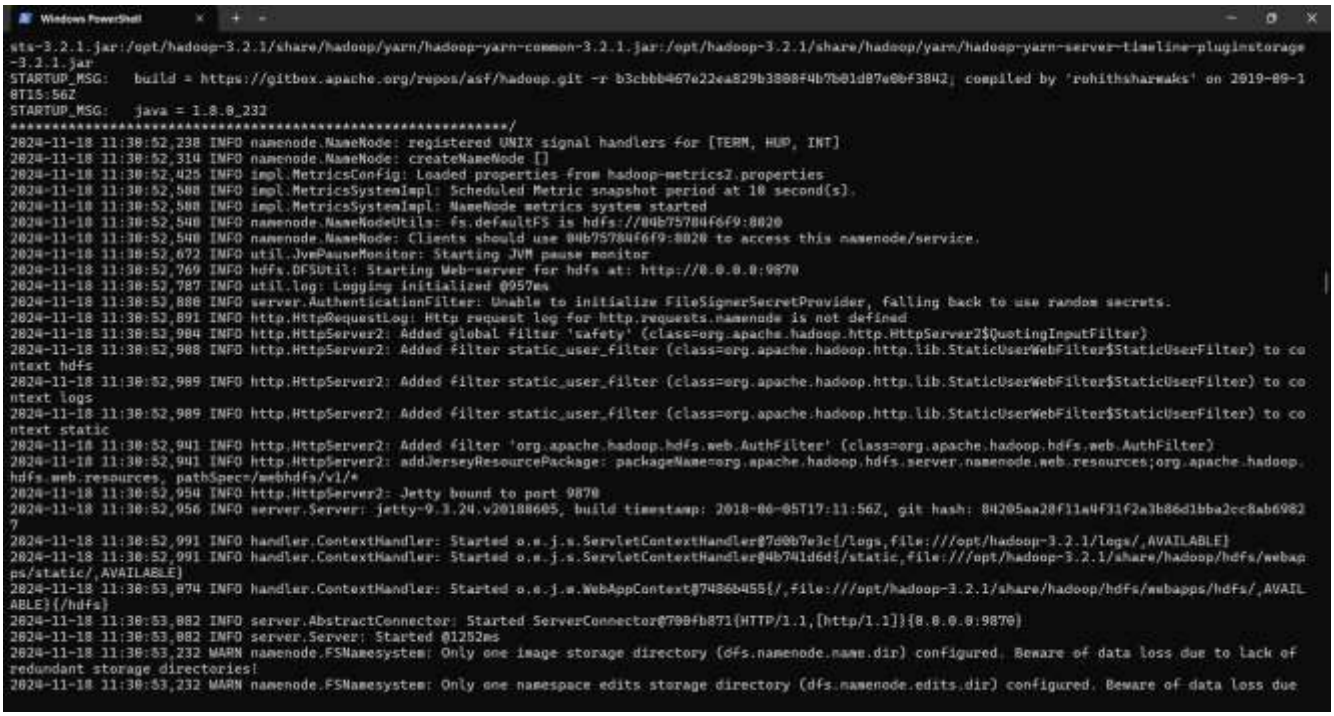
Windows PowerShell
sts-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-common-3.2.1.jar:/opt/hadoop-3.2.1/share/hadoop/yarn/hadoop-yarn-server-timeline-pluginstorage
-3.2.1.jar
STARTUP_MSG: build = https://gitbox.apache.org/repos/asf/hadoop.git -r b3cbbb467e22ea829b3808f4b7b01d07e0bf3842; compiled by 'rohithsharmaks' on 2019-09-1
8T15:56Z
STARTUP_MSG: java = 1.8.0_232
**/
2024-11-18 11:30:52,238 INFO namenode.NameNode: registered UNIX signal handlers for [TERM, HUP, INT]
2024-11-18 11:30:52,314 INFO namenode.NameNode: createNameNode []
2024-11-18 11:30:52,425 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2024-11-18 11:30:52,508 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2024-11-18 11:30:52,508 INFO impl.MetricsSystemImpl: NameNode metrics system started
2024-11-18 11:30:52,540 INFO namenode.NameNodeUtils: fs.defaultFS is hdfs://04b75784f6f9:8020
2024-11-18 11:30:52,540 INFO namenode.NameNode: Clients should use 04b75784f6f9:8020 to access this namenode/service.
2024-11-18 11:30:52,672 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2024-11-18 11:30:52,769 INFO hdfs.DFSUtil: Starting Web-server for hdfs at: http://0.0.0.0:9870
2024-11-18 11:30:52,787 INFO util.log: Logging initialized @957ms
2024-11-18 11:30:53,880 INFO server.AuthenticationFilter: Unable to initialize FileSignerSecretProvider, falling back to use random secrets.
2024-11-18 11:30:53,891 INFO http.HttpRequestLog: Http request log for http.requests.namenode is not defined
2024-11-18 11:30:53,904 INFO http.HttpServer2: Added global filter 'safety' (class=org.apache.hadoop.http.HttpServer2$QuotingInputFilter)
2024-11-18 11:30:53,908 INFO http.HttpServer2: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to co
ntext hdfs
2024-11-18 11:30:53,909 INFO http.HttpServer2: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to co
ntext logs
2024-11-18 11:30:53,909 INFO http.HttpServer2: Added filter static_user_filter (class=org.apache.hadoop.http.lib.StaticUserWebFilter$StaticUserFilter) to co
ntext static
2024-11-18 11:30:53,941 INFO http.HttpServer2: Added filter 'org.apache.hadoop.hdfs.web.AuthFilter' (class=org.apache.hadoop.hdfs.web.AuthFilter)
2024-11-18 11:30:53,941 INFO http.HttpServer2: addJerseyResourcePackage: packageName=org.apache.hadoop.hdfs.server.namenode.web.resources;org.apache.hadoop.
hdfs.web.resources, pathSpec=/webhdfs/v1/*
2024-11-18 11:30:53,954 INFO http.HttpServer2: Jetty bound to port 9870
2024-11-18 11:30:53,956 INFO server.Server: jetty-9.3.24.v20180605, build timestamp: 2018-06-05T17:11:56Z, git hash: 84205aa28f11a4f31f2a3b86d1bba2cc8ab6982
7
2024-11-18 11:30:53,991 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@7d0b7e3c{/logs,file:///opt/hadoop-3.2.1/logs/,AVAILABLE}
2024-11-18 11:30:53,991 INFO handler.ContextHandler: Started o.e.j.s.ServletContextHandler@4b741d6d{/static,file:///opt/hadoop-3.2.1/share/hadoop/hdfs/webap
ps/static/,AVAILABLE}
2024-11-18 11:30:53,874 INFO handler.ContextHandler: Started o.e.j.w.WebAppContext@74866455{/,file:///opt/hadoop-3.2.1/share/hadoop/hdfs/webapps/hdfs/,AVAIL
ABLE}{/hdfs}
2024-11-18 11:30:53,882 INFO server.AbstractConnector: Started ServerConnector@700fb871{HTTP/1.1,[http/1.1]}{0.0.0.0:9870}
2024-11-18 11:30:53,882 INFO server.Server: Started @1252ms
2024-11-18 11:30:53,232 WARN namenode.FSNamesystem: Only one image storage directory (dfs.namenode.name.dir) configured. Beware of data loss due to lack of
redundant storage directories!
2024-11-18 11:30:53,232 WARN namenode.FSNamesystem: Only one namespace edits storage directory (dfs.namenode.edits.dir) configured. Beware of data loss due

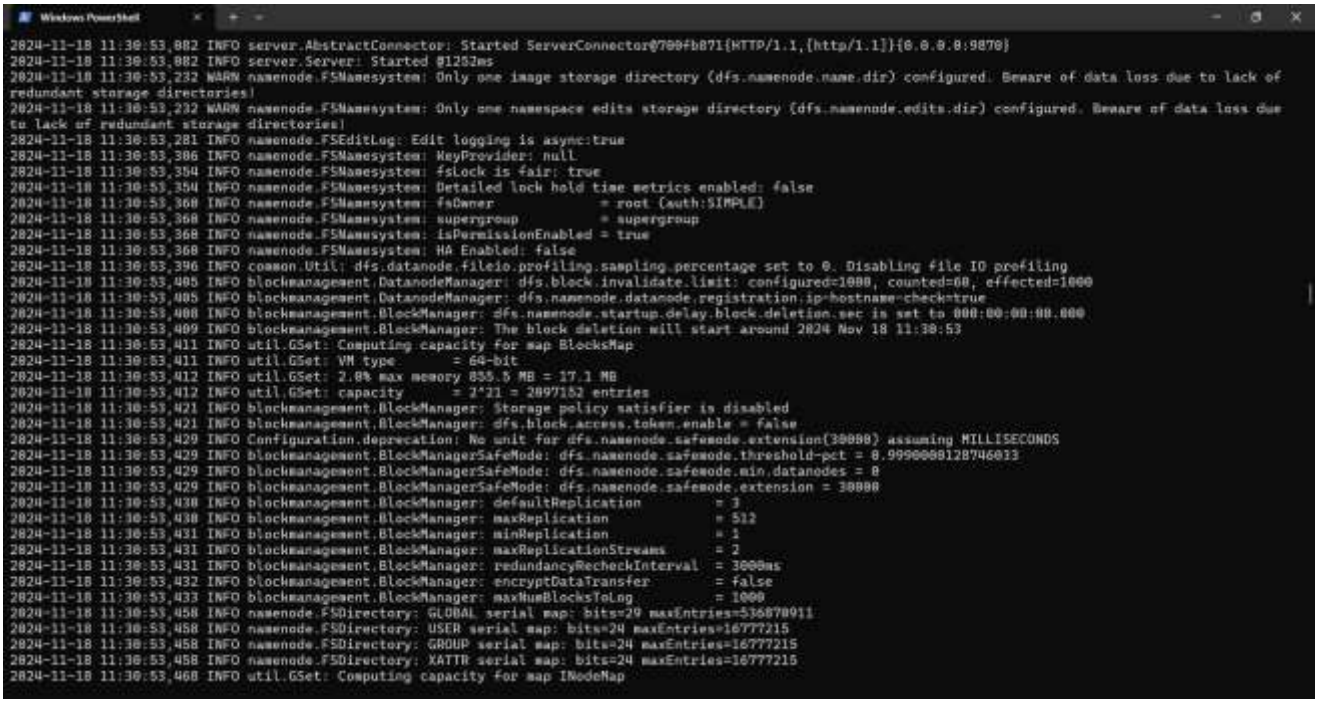

```
2024-11-18 11:30:53,882 INFO server.AbstractConnector: Started ServerConnector@700fb871{HTTP/1.1,[http/1.1]}{0.0.0.0:9870}
2024-11-18 11:30:53,882 INFO server.Server: Started @1252ms
2024-11-18 11:30:53,232 WARN namenode.FSNamesystem: Only one image storage directory (dfs.namenode.name.dir) configured. Beware of data loss due to lack of
redundant storage directories!
2024-11-18 11:30:53,232 WARN namenode.FSNamesystem: Only one namespace edits storage directory (dfs.namenode.edits.dir) configured. Beware of data loss due
to lack of redundant storage directories!
2024-11-18 11:30:53,281 INFO namenode.FSEditLog: Edit logging is async:true
2024-11-18 11:30:53,306 INFO namenode.FSNamesystem: KeyProvider: null
2024-11-18 11:30:53,354 INFO namenode.FSNamesystem: fsLock is fair: true
2024-11-18 11:30:53,354 INFO namenode.FSNamesystem: Detailed lock hold time metrics enabled: false
2024-11-18 11:30:53,360 INFO namenode.FSNamesystem: fsOwner             = root (auth:SIMPLE)
2024-11-18 11:30:53,360 INFO namenode.FSNamesystem: supergroup          = supergroup
2024-11-18 11:30:53,360 INFO namenode.FSNamesystem: isPermissionEnabled = true
2024-11-18 11:30:53,360 INFO namenode.FSNamesystem: HA Enabled: false
2024-11-18 11:30:53,396 INFO common.Util: dfs.datanode.fileio.profiling.sampling.percentage set to 0. Disabling file IO profiling
2024-11-18 11:30:53,405 INFO blockmanagement.DatanodeManager: dfs.block.invalidate.limit: configured=1000, counted=60, effected=1000
2024-11-18 11:30:53,405 INFO blockmanagement.DatanodeManager: dfs.namenode.datanode.registration.ip-hostname-check=true
2024-11-18 11:30:53,408 INFO blockmanagement.BlockManager: dfs.namenode.startup.delay.block.deletion.sec is set to 000:00:00:00.000
2024-11-18 11:30:53,409 INFO blockmanagement.BlockManager: The block deletion will start around 2024 Nov 18 11:30:53
2024-11-18 11:30:53,411 INFO util.GSet: Computing capacity for map BlocksMap
2024-11-18 11:30:53,411 INFO util.GSet: VM type       = 64-bit
2024-11-18 11:30:53,412 INFO util.GSet: 2.0% max memory 855.5 MB = 17.1 MB
2024-11-18 11:30:53,412 INFO util.GSet: capacity      = 2^21 = 2097152 entries
2024-11-18 11:30:53,421 INFO blockmanagement.BlockManager: Storage policy satisfier is disabled
2024-11-18 11:30:53,421 INFO blockmanagement.BlockManager: dfs.block.access.token.enable = false
2024-11-18 11:30:53,429 INFO Configuration.deprecation: No unit for dfs.namenode.safemode.extension(30000) assuming MILLISECONDS
2024-11-18 11:30:53,429 INFO blockmanagement.BlockManagerSafeMode: dfs.namenode.safemode.threshold-pct = 0.9990000128746033
2024-11-18 11:30:53,429 INFO blockmanagement.BlockManagerSafeMode: dfs.namenode.safemode.min.datanodes = 0
2024-11-18 11:30:53,429 INFO blockmanagement.BlockManagerSafeMode: dfs.namenode.safemode.extension = 30000
2024-11-18 11:30:53,430 INFO blockmanagement.BlockManager: defaultReplication         = 3
2024-11-18 11:30:53,430 INFO blockmanagement.BlockManager: maxReplication             = 512
2024-11-18 11:30:53,431 INFO blockmanagement.BlockManager: minReplication             = 1
2024-11-18 11:30:53,431 INFO blockmanagement.BlockManager: maxReplicationStreams      = 2
2024-11-18 11:30:53,431 INFO blockmanagement.BlockManager: redundancyRecheckInterval  = 3000ms
2024-11-18 11:30:53,432 INFO blockmanagement.BlockManager: encryptDataTransfer        = false
2024-11-18 11:30:53,433 INFO blockmanagement.BlockManager: maxNumBlocksToLog          = 1000
2024-11-18 11:30:53,458 INFO namenode.FSDirectory: GLOBAL serial map: bits=29 maxEntries=536870911
2024-11-18 11:30:53,458 INFO namenode.FSDirectory: USER serial map: bits=24 maxEntries=16777215
2024-11-18 11:30:53,458 INFO namenode.FSDirectory: GROUP serial map: bits=24 maxEntries=16777215
2024-11-18 11:30:53,458 INFO namenode.FSDirectory: XATTR serial map: bits=24 maxEntries=16777215
2024-11-18 11:30:53,468 INFO util.GSet: Computing capacity for map INodeMap
```

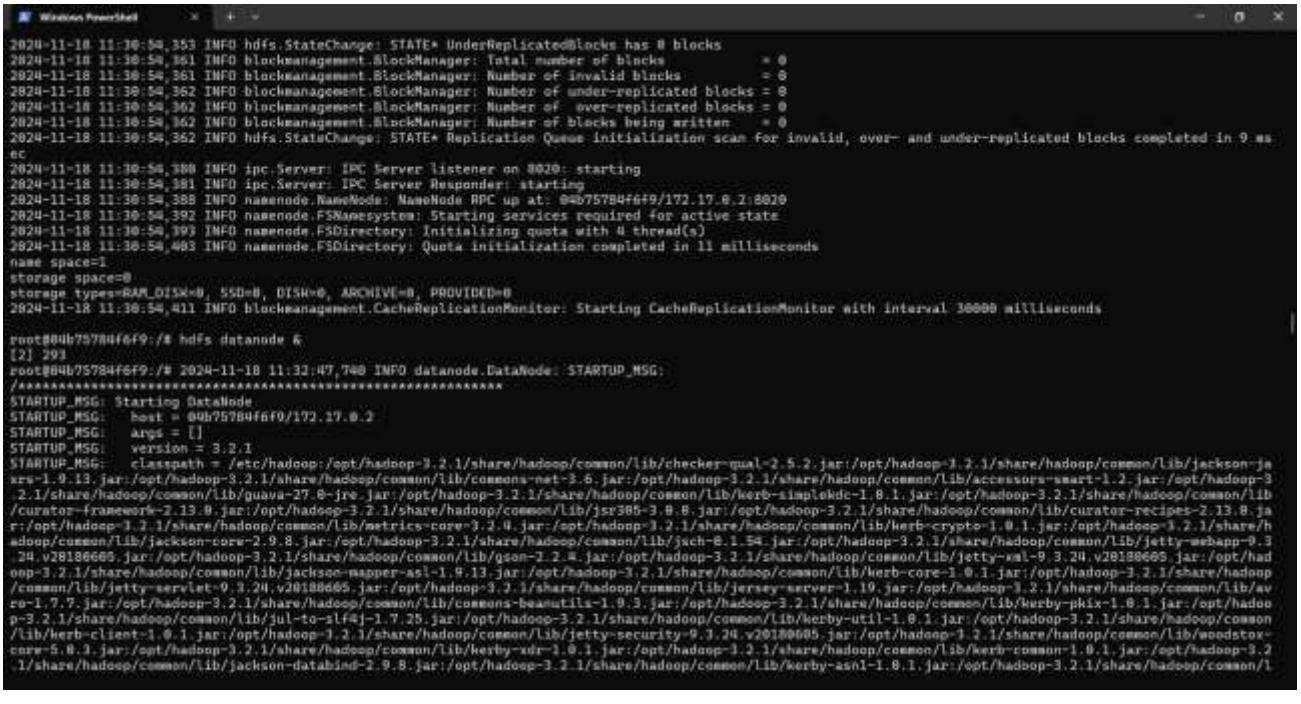

```
2024-11-18 11:30:54,353 INFO hdfs.StateChange: STATE* UnderReplicatedBlocks has 0 blocks
2024-11-18 11:30:54,361 INFO blockmanagement.BlockManager: Total number of blocks            = 0
2024-11-18 11:30:54,361 INFO blockmanagement.BlockManager: Number of invalid blocks          = 0
2024-11-18 11:30:54,362 INFO blockmanagement.BlockManager: Number of under-replicated blocks = 0
2024-11-18 11:30:54,362 INFO blockmanagement.BlockManager: Number of  over-replicated blocks = 0
2024-11-18 11:30:54,362 INFO blockmanagement.BlockManager: Number of blocks being written    = 0
2024-11-18 11:30:54,362 INFO hdfs.StateChange: STATE* Replication Queue initialization scan for invalid, over- and under-replicated blocks completed in 9 ms
ec
2024-11-18 11:30:54,380 INFO ipc.Server: IPC Server listener on 8020: starting
2024-11-18 11:30:54,381 INFO ipc.Server: IPC Server Responder: starting
2024-11-18 11:30:54,388 INFO namenode.NameNode: NameNode RPC up at: 04b75784f6f9/172.17.0.2:8020
2024-11-18 11:30:54,392 INFO namenode.FSNamesystem: Starting services required for active state
2024-11-18 11:30:54,393 INFO namenode.FSDirectory: Initializing quota with 4 thread(s)
2024-11-18 11:30:54,403 INFO namenode.FSDirectory: Quota initialization completed in 11 milliseconds
name space=1
storage space=0
storage types=RAM_DISK=0, SSD=0, DISK=0, ARCHIVE=0, PROVIDED=0
2024-11-18 11:30:54,411 INFO blockmanagement.CacheReplicationMonitor: Starting CacheReplicationMonitor with interval 30000 milliseconds

root@04b75784f6f9:/# hdfs datanode &
[2] 293
root@04b75784f6f9:/# 2024-11-18 11:32:47,740 INFO datanode.DataNode: STARTUP_MSG:
/************************************************************
STARTUP_MSG: Starting DataNode
STARTUP_MSG:   host = 04b75784f6f9/172.17.0.2
STARTUP_MSG:   args = []
STARTUP_MSG:   version = 3.2.1
STARTUP_MSG:   classpath = /etc/hadoop:/opt/hadoop-3.2.1/share/hadoop/common/lib/checker-qual-2.5.2.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jackson-ja
xrs-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-net-3.6.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/accessors-smart-1.2.jar:/opt/hadoop-3
.2.1/share/hadoop/common/lib/guava-27.0-jre.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-simplekdc-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib
/curator-framework-2.13.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jsr305-3.0.0.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/curator-recipes-2.13.0.ja
r:/opt/hadoop-3.2.1/share/hadoop/common/lib/metrics-core-3.2.4.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-crypto-1.0.1.jar:/opt/hadoop-3.2.1/share/h
adoop/common/lib/jackson-core-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jsch-0.1.54.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-webapp-9.3
.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/gson-2.2.4.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-xml-9.3.24.v20180605.jar:/opt/had
oop-3.2.1/share/hadoop/common/lib/jackson-mapper-asl-1.9.13.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-core-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop
/common/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jersey-server-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/av
ro-1.7.7.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/commons-beanutils-1.9.3.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-pkix-1.0.1.jar:/opt/hadoo
p-3.2.1/share/hadoop/common/lib/jul-to-slf4j-1.7.25.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-util-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common
/lib/kerb-client-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-security-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/woodstox-
core-5.0.3.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-xdr-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerb-common-1.0.1.jar:/opt/hadoop-3.2
.1/share/hadoop/common/lib/jackson-databind-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/kerby-asn1-1.0.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/l
```

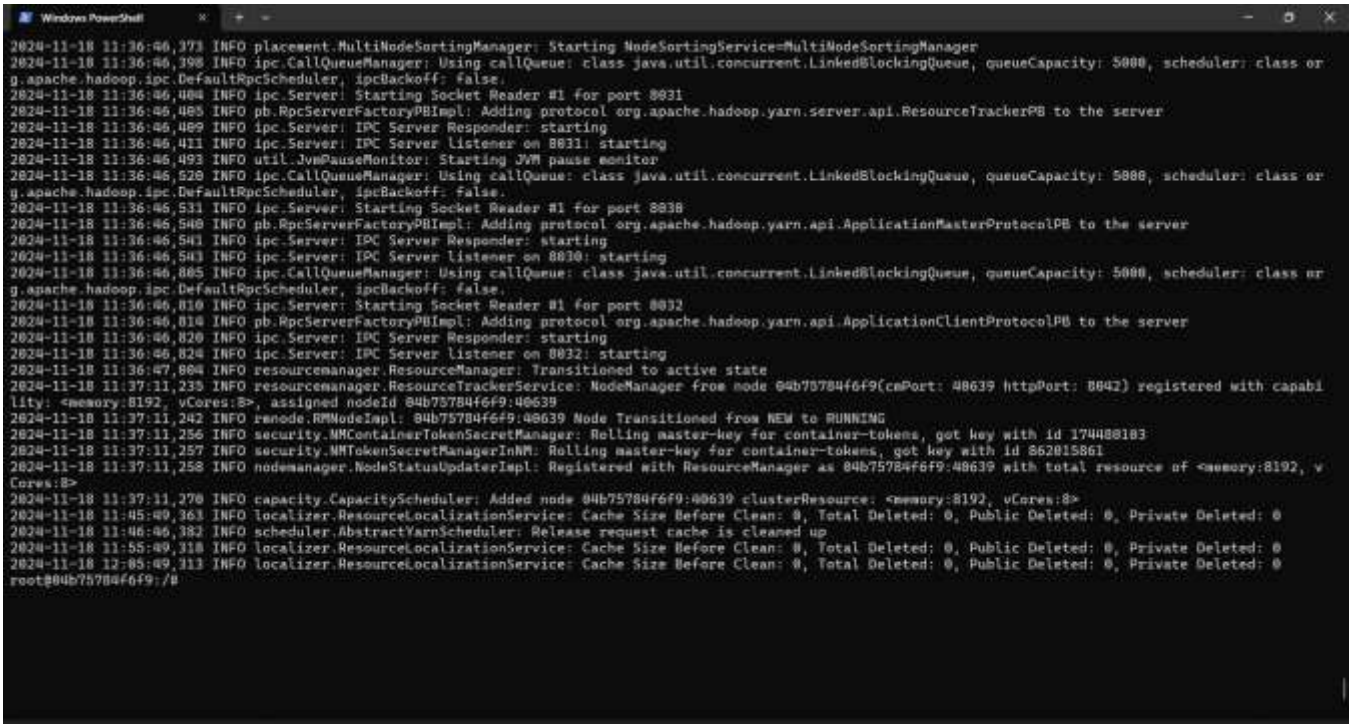
Windows PowerShell
2024-11-18 11:36:46,373 INFO placement.MultiNodeSortingManager: Starting NodeSortingService=MultiNodeSortingManager
2024-11-18 11:36:46,398 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class or
g.apache.hadoop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2024-11-18 11:36:46,404 INFO ipc.Server: Starting Socket Reader #1 for port 8031
2024-11-18 11:36:46,405 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.hadoop.yarn.server.api.ResourceTrackerPB to the server
2024-11-18 11:36:46,409 INFO ipc.Server: IPC Server Responder: starting
2024-11-18 11:36:46,411 INFO ipc.Server: IPC Server listener on 8031: starting
2024-11-18 11:36:46,493 INFO util.JvmPauseMonitor: Starting JVM pause monitor
2024-11-18 11:36:46,520 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class or
g.apache.hadoop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2024-11-18 11:36:46,531 INFO ipc.Server: Starting Socket Reader #1 for port 8030
2024-11-18 11:36:46,540 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.hadoop.yarn.api.ApplicationMasterProtocolPB to the server
2024-11-18 11:36:46,541 INFO ipc.Server: IPC Server Responder: starting
2024-11-18 11:36:46,543 INFO ipc.Server: IPC Server listener on 8030: starting
2024-11-18 11:36:46,805 INFO ipc.CallQueueManager: Using callQueue: class java.util.concurrent.LinkedBlockingQueue, queueCapacity: 5000, scheduler: class or
g.apache.hadoop.ipc.DefaultRpcScheduler, ipcBackoff: false.
2024-11-18 11:36:46,810 INFO ipc.Server: Starting Socket Reader #1 for port 8032
2024-11-18 11:36:46,814 INFO pb.RpcServerFactoryPBImpl: Adding protocol org.apache.hadoop.yarn.api.ApplicationClientProtocolPB to the server
2024-11-18 11:36:46,820 INFO ipc.Server: IPC Server Responder: starting
2024-11-18 11:36:46,824 INFO ipc.Server: IPC Server listener on 8032: starting
2024-11-18 11:36:47,804 INFO resourcemanager.ResourceManager: Transitioned to active state
2024-11-18 11:37:11,235 INFO resourcemanager.ResourceTrackerService: NodeManager from node 04b75784f6f9(cmPort: 40639 httpPort: 8042) registered with capabi
lity: <memory:8192, vCores:8>, assigned nodeId 04b75784f6f9:40639
2024-11-18 11:37:11,242 INFO rmnode.RMNodeImpl: 04b75784f6f9:40639 Node Transitioned from NEW to RUNNING
2024-11-18 11:37:11,256 INFO security.NMContainerTokenSecretManager: Rolling master-key for container-tokens, got key with id 174488183
2024-11-18 11:37:11,257 INFO security.NMTokenSecretManagerInNM: Rolling master-key for container-tokens, got key with id 862015861
2024-11-18 11:37:11,258 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as 04b75784f6f9:40639 with total resource of <memory:8192, v
Cores:8>
2024-11-18 11:37:11,270 INFO capacity.CapacityScheduler: Added node 04b75784f6f9:40639 clusterResource: <memory:8192, vCores:8>
2024-11-18 11:45:49,363 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
2024-11-18 11:46:46,382 INFO scheduler.AbstractYarnScheduler: Release request cache is cleaned up
2024-11-18 11:55:49,318 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
2024-11-18 12:05:49,313 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public Deleted: 0, Private Deleted: 0
root@04b75784f6f9:/#